

Concurrency Control: 2PL

Dr. L.M. Jenila Livingston
VIT Chennai

Lock Management

- A lock is a mechanism to **control concurrent access** to a data item
- Lock requests are made to **concurrency-control manager** or **lock manager**. Transaction can proceed only after request is granted.
- **Lock table** entry:
 - **Number of transactions** currently holding a lock
 - **Type of lock** held (shared or exclusive)
 - **Pointer to queue** of lock requests

3 Types of Locks

- We allow transactions to **lock** objects.

Shared lock (S): Data item can **only be read**. S-lock is requested using **lock-S** instruction.

Exclusive lock (X): Data item can be **both read** as well as **write**. X-lock is requested using **lock-X** instruction.

Lock-Based Protocols (Cont.)

■ Lock-compatibility matrix


	S	X
S	true	false
X	false	false

- **Any number of transactions can hold shared locks on an item,**
 - but if any transaction holds an exclusive on the item no other transaction may hold any lock on the item.
- If a lock cannot be granted, the requesting transaction is made to wait till all incompatible locks held by other transactions have been released. The lock is then granted.

The Two-Phase Locking (2PL) Protocol

- This is a protocol which ensures conflict-serializable schedules.
- Phase 1: Growing Phase
 - transaction may obtain locks
 - transaction may not release locks
- Phase 2: Shrinking Phase
 - transaction may release locks
 - transaction may not obtain locks
- The protocol assures serializability. It can be proved that the transactions can be serialized in the order of their **lock points** (i.e. the point where a transaction acquired its final lock).

2 PL



The Two-Phase Locking Protocol (Cont.)

- A **locking protocol** is a set of rules followed by all transactions while requesting and releasing locks. Locking protocols restrict the set of possible schedules
- Two-phase locking *does not* ensure freedom from deadlocks
- **Cascading roll-back is possible** under two-phase locking. To avoid this, follow a modified protocol called **strict two-phase locking**. Here a transaction must hold all its **exclusive locks** till it commits/aborts.
- **Rigorous two-phase locking** is even stricter: here **all locks** are held till commit/abort. In this protocol transactions can be serialized in the order in which they commit.
- .

Lock Conversions

- Two-phase locking with lock conversions:
 - First Phase:
 - can acquire a lock-S on item
 - can acquire a lock-X on item
 - can convert a lock-S to a lock-X (upgrade)
 - Second Phase:
 - can release a lock-S
 - can release a lock-X
 - can convert a lock-X to a lock-S (downgrade)
- This protocol assures serializability. But still relies on the programmer to insert the various locking instructions.

2 PL: Automatic Acquisition of Locks

- read

- A transaction T_i issues the standard read/write instruction, without explicit locking calls.
- The operation **read**(D) is processed as:

```
if  $T_i$  has a lock on  $D$  // lock-X OR lock-S
then
    read( $D$ )
else begin
    if necessary wait until no other
        transaction has a lock-X on  $D$ 
    grant  $T_i$  a lock-S on  $D$ ;
    read( $D$ )
end
```

2PL: Automatic Acquisition of Locks

- write

- **write(D)** is processed as:

if T_i has a **lock-X** on D

then

 write(D)

else begin

 if necessary wait until no other trans. has any lock on D ,

 if T_i has a **lock-S** on D

then

upgrade lock on D to **lock-X**

else

 grant T_i a **lock-X** on D

 write(D)

end;

- **All locks are released after commit or abort**

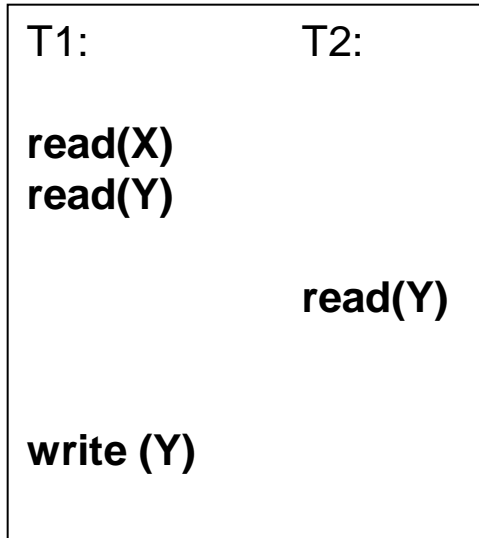
Lock-Based Protocols (Cont.)

- **Example** of a transaction performing locking:

T_2 : **begin**
 lock-S(A);
 read (A);
 unlock(A);
 lock-S(B);
 read (B);
 unlock(B);
 display(A+B)
 commit;

<p>T2: read(A) read(B) display(A+B)</p>
--

Lock-Based Protocols – 2PL



T1:

begin
lock-S(X)
read(X)
lock-S(Y)
read(Y)

upgrade(Y)
write (Y)
unlock(X)
unlock(Y)
commit

T2:

begin
lock-S(Y)
read(Y)
unlock(Y)
commit

Pitfalls of Lock-Based Protocols

- Consider the partial schedule

T_3	T_4
lock-x (B)	
read (B)	
$B := B - 50$	
write (B)	
	lock-s (A)
	read (A)
	lock-s (B)
lock-x (A)	

- Neither T_3 nor T_4 can make progress — executing **lock-S(B)** causes T_4 to wait for T_3 to release its lock on B , while executing **lock-X(A)** causes T_3 to wait for T_4 to release its lock on A .
- Such a situation is called a **deadlock**.
 - To handle a deadlock one of T_3 or T_4 must be rolled back and its locks released.


Pitfalls of Lock-Based Protocols (Cont.)

- **Starvation** is also possible if concurrency control manager is badly designed. For example:
 - A transaction may be waiting for an X-lock on an item, while a sequence of other transactions request and are granted an S-lock on the same item.
 - The same transaction is repeatedly rolled back due to deadlocks.
- Concurrency control manager can be designed to prevent starvation.

Implementation of Locking

- A **lock manager** can be implemented as a separate process to which transactions send lock and unlock requests
- The lock manager replies to a lock request by sending a lock grant messages (or a message asking the transaction to roll back, in case of a deadlock)
- The requesting transaction waits until its request is answered
- The **lock manager maintains a data-structure** called a **lock table** to record granted locks and pending requests
- The lock table is usually implemented as an in-memory hash table indexed on the name of the data item being locked

Lock Table



- Lock table records the type of lock granted or requested
- New request is added to the end of the queue of requests for the data item, and granted if it is compatible with all earlier locks
- Unlock requests result in the request being deleted, and later requests are checked to see if they can now be granted
- If transaction aborts, all waiting or granted requests of the transaction are deleted
 - lock manager may keep a list of locks held by each transaction, to implement this efficiently

Timestamp-Based Protocols

1. Use the value of the system clock as the timestamp; that is, a transaction's timestamp is equal to the value of the clock when the transaction enters the system.
2. Use a logical counter that is incremented after a new timestamp has been assigned; that is, a transaction's timestamp is equal to the value of the counter when the transaction enters the system.

Timestamp-based Protocols

- Suppose there are there transactions T1, T2, and T3.
- T1 has entered the system at time 0010
- T2 has entered the system at 0020
- T3 has entered the system at 0030
- Priority will be given to transaction T1, then transaction T2 and lastly Transaction T3.

Thank You!